

```
In [93]: #Hlanhla Hlungwane
#Python project on Jupyter notebook
#Employee Turnover Analytics
#09 July 2025
```

```
In [94]: import numpy as np    #Allows computations in Linear Algebra
import pandas as pd    #Allows for data processing
import os
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [95]: df = pd.read_csv('HR_comma_sep.csv')
```

```
In [96]: df.head()
```

```
Out[96]:
```

| | satisfaction_level | last_evaluation | number_project | average_monthly_hours | time_spent_cc |
|---|--------------------|-----------------|----------------|-----------------------|---------------|
| 0 | 0.38 | 0.53 | 2 | 157 | |
| 1 | 0.80 | 0.86 | 5 | 262 | |
| 2 | 0.11 | 0.88 | 7 | 272 | |
| 3 | 0.72 | 0.87 | 5 | 223 | |
| 4 | 0.37 | 0.52 | 2 | 159 | |

```
In [97]: df.describe()
```

```
Out[97]:
```

| | satisfaction_level | last_evaluation | number_project | average_monthly_hours | time_spent |
|-------|--------------------|-----------------|----------------|-----------------------|------------|
| count | 14999.000000 | 14999.000000 | 14999.000000 | 14999.000000 | 14 |
| mean | 0.612834 | 0.716102 | 3.803054 | 201.050337 | |
| std | 0.248631 | 0.171169 | 1.232592 | 49.943099 | |
| min | 0.090000 | 0.360000 | 2.000000 | 96.000000 | |
| 25% | 0.440000 | 0.560000 | 3.000000 | 156.000000 | |
| 50% | 0.640000 | 0.720000 | 4.000000 | 200.000000 | |
| 75% | 0.820000 | 0.870000 | 5.000000 | 245.000000 | |
| max | 1.000000 | 1.000000 | 7.000000 | 310.000000 | |

```
In [98]: #Train and Test Data
from sklearn.model_selection import train_test_split
X = df.drop(columns = ['left'])
y = df.left
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 1/3)
```

```
In [99]: print('Ratio of workers who left in train:', y_train.sum()/y_train.count())
Ratio of workers who left in train: 0.23522352235223523
```


```
In [100]: print('Ratio of workers who left in test:', y_test.sum()/y_test.count())
Ratio of workers who left in test: 0.2438
```

```
In [101... print('Accuracy:', 1-y_test.sum()/y_test.size)
```

Accuracy: 0.7562

```
In [102... #Data Cleaning  
sns.barplot(x=X_train.salary, y=y_train, order = ['low', 'medium', 'high'])
```

Out[102]: <AxesSubplot: xlabel='salary', ylabel='left'>



```
In [103... #Employees with low salaries left the company the most  
#The employees earning higher are less likely leave the company
```

```
In [104... #The label encoder will be labelled with low, medium and high salaries  
#These will be converted to 0, 1, 2
```

```
In [105... #Encoding is when categorical data is converted to numerical data
```

```
In [106... salary_encoder = {'low': 0, 'medium': 1, 'higher':2}  
X_train['salary'] = X_train.salary.map(salary_encoder)  
X_test['salary'] = X_test.salary.map(salary_encoder)
```

```
In [107... from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import make_column_transformer

dept_encoder = OneHotEncoder(handle_unknown='ignore')
dept_encoder.fit_transform(X_train[['sales']])
dept_encoder.transform(X_test[['sales']])

transformer = make_column_transformer(
    (OneHotEncoder(), ['sales']),
    remainder='passthrough')

transformed = transformer.fit_transform(X_train)
X_train = pd.DataFrame(
    transformed,
    columns=transformer.get_feature_names_out()
)

transformed = transformer.transform(X_test)
X_test = pd.DataFrame(
    transformed,
    columns=transformer.get_feature_names_out()
)
```

```
In [108... X = X_train.copy()
X['left'] = y_train
```

```
In [109... #Visualizations

sns.pairplot(data=df, hue = 'left')
```

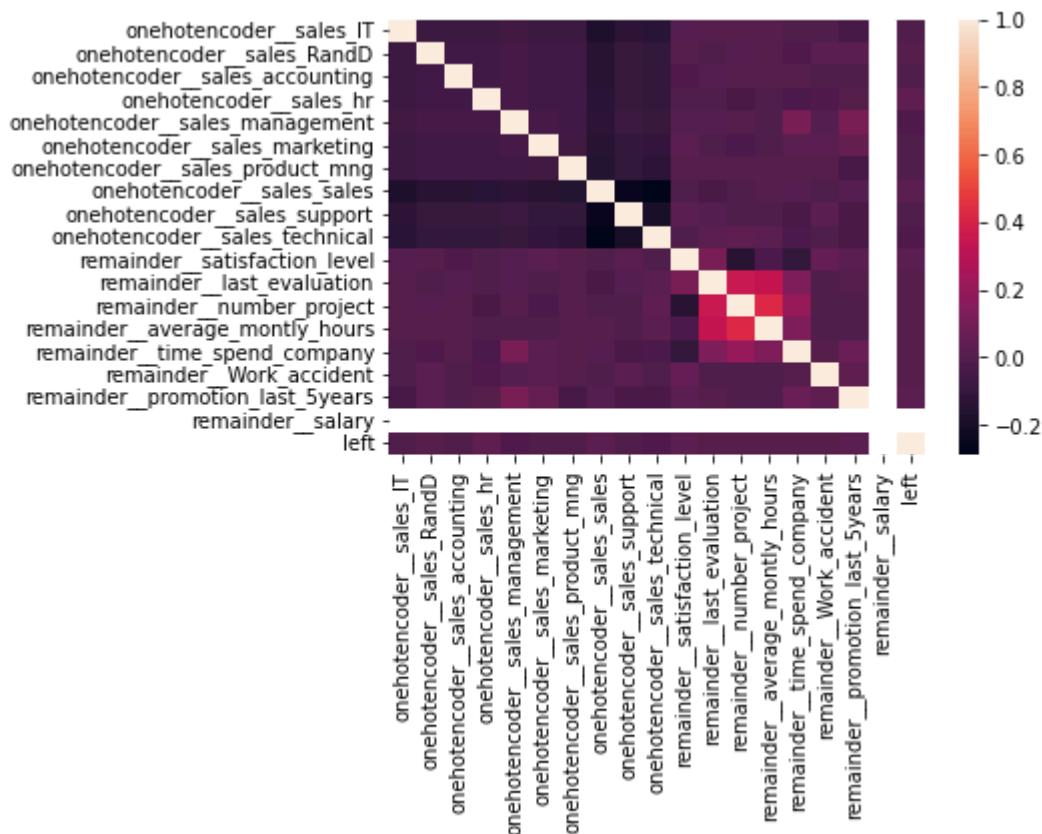
```
Out[109]: <seaborn.axisgrid.PairGrid at 0x7ffa4f69de40>
```



In [110]: `#Heat map plot`

```
sns.heatmap(X.corr(),
             xticklabels = X.columns,
             yticklabels = X.columns)
```

Out[110]: `<AxesSubplot: >`



```
In [111... from sklearn.linear_model import RidgeClassifierCV
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer # Import for handling missing values
import random

# Define transformers
numeric_transformer = Pipeline(
    steps=[
        ('imputer', SimpleImputer(strategy='mean')), # Handle missing values
        ('scaler', StandardScaler()) # Scale numeric features
    ]
)

categorical_transformer = Pipeline(
    steps=[
        ('imputer', SimpleImputer(strategy='most_frequent')), # Handle missing values
        ('onehot', OneHotEncoder(handle_unknown='ignore')) # Encode categorical features
    ]
)

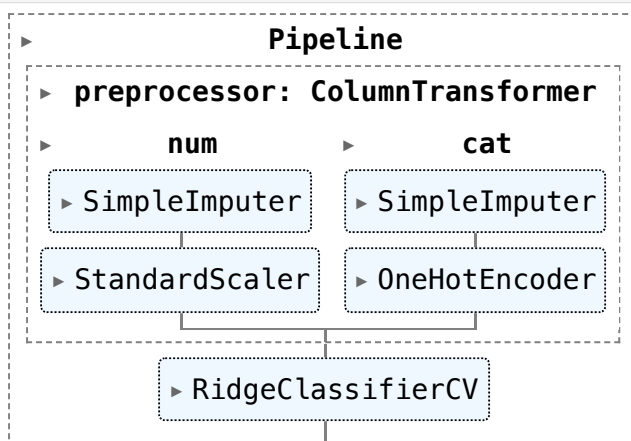
# Combine transformers into a preprocessor
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, X_train.select_dtypes(include=['float64', 'int64'])),
        ('cat', categorical_transformer, X_train.select_dtypes(include=['object']))
    ]
)

# Define the model pipeline
model = Pipeline(
    steps=[
        ('preprocessor', preprocessor),
        ('model', RidgeClassifierCV(alphas=[random.uniform(0.1, 1) * 10**x for x in range(0, 10)]))
    ]
)
```

```
]
)

# Fit the model
model.fit(X_train, y_train)
```

Out[111]:



```
In [112... model.fit(X_train,y=y_train)
model['model'].best_score_
```

Out[112]: 0.7771770885442721

```
In [113... #A stronger model should be built
df.head(2)
```

```
Out[113]:
```

| | satisfaction_level | last_evaluation | number_project | average_monthly_hours | time_spend_c |
|---|--------------------|-----------------|----------------|-----------------------|--------------|
| 0 | 0.38 | 0.53 | 2 | 157 | |
| 1 | 0.80 | 0.86 | 5 | 262 | |

```
In [114... print(X_train.columns)

Index(['onehotencoder__sales_IT', 'onehotencoder__sales_RandD',
      'onehotencoder__sales_accounting', 'onehotencoder__sales_hr',
      'onehotencoder__sales_management', 'onehotencoder__sales_marketing',
      'onehotencoder__sales_product_mng', 'onehotencoder__sales_sales',
      'onehotencoder__sales_support', 'onehotencoder__sales_technical',
      'remainder__satisfaction_level', 'remainder__last_evaluation',
      'remainder__number_project', 'remainder__average_monthly_hours',
      'remainder__time_spend_company', 'remainder__Work_accident',
      'remainder__promotion_last_5years', 'remainder__salary'],
      dtype='object')
```

```
In [125... from sklearn.preprocessing import PolynomialFeatures, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import RidgeClassifierCV
import random

# Define the preprocessor
preprocessor = ColumnTransformer(
    transformers=[
        # Use the prefixed column names for PolynomialFeatures
        ("poly", PolynomialFeatures(8), [
            'remainder__satisfaction_level',
            'remainder__last_evaluation',
            'remainder__average_monthly_hours'
```

```

    ]),
    ("poly2", PolynomialFeatures(2), [
        'remainder__number_project',
        'remainder__time_spend_company'
    ]),
    # Apply StandardScaler to all columns in X_train
    ('scaler', StandardScaler(), X_train.columns)
]
)

# Define the pipeline
model = Pipeline(
    steps=[
        ("preprocessor", preprocessor),
        ("model", RidgeClassifierCV(alphas=[random.uniform(0.1, 1) * 10**x for
    ]
)

# Fit the model
model.fit(X_train, y_train)

# Access the best score
print(model['model'].best_score_)

0.9458947973986993

```

```

In [126... print('Accuracy on test data:', model.score(X_test,y_test))

Accuracy on test data: 0.9422

```

```

In [127... from xgboost.sklearn import XGBClassifier
from sklearn.model_selection import GridSearchCV

#The model consists of the scaler and XGBClassifier
clf = Pipeline(
    steps=[('scaler', StandardScaler()),
        ("xgb", XGBClassifier(cv=5))
    ]
)

param_grid = {
    "xgb__n_estimators": [300],
    "xgb__learning_rate": [0.3],
    "xgb__eta": [0.1],
    "xgb__max_depth": [1],
    "xgb__objective": ['reg:squarederror'],
    'xgb__min_child_weight': [4],
    'xgb__subsample': [0.7],
    'xgb__reg_lambda': [0.55]
}

search = GridSearchCV(clf, param_grid=param_grid, cv = 5, scoring='accuracy')

search.fit(X_train,y_train)
print(search.best_params_)

```

```
[08:48:34] WARNING: ../src/learner.cc:627:  
Parameters: { "cv" } might not be used.
```

This could be a false alarm, with some parameters getting used by language bindings but then being mistakenly passed down to XGBoost core, or some parameter actually being used but getting flagged wrongly here. Please open an issue if you find any such cases.

```
{'xgb__eta': 0.1, 'xgb__learning_rate': 0.3, 'xgb__max_depth': 1, 'xgb__min  
_child_weight': 4, 'xgb__n_estimators': 300, 'xgb__objective': 'reg:squared  
error', 'xgb__reg_lambda': 0.55, 'xgb__subsample': 0.7}
```

```
In [128... search.best_score_
```

```
Out[128]: 0.9491942471235617
```

```
In [ ]: #The model has been strengthened
```